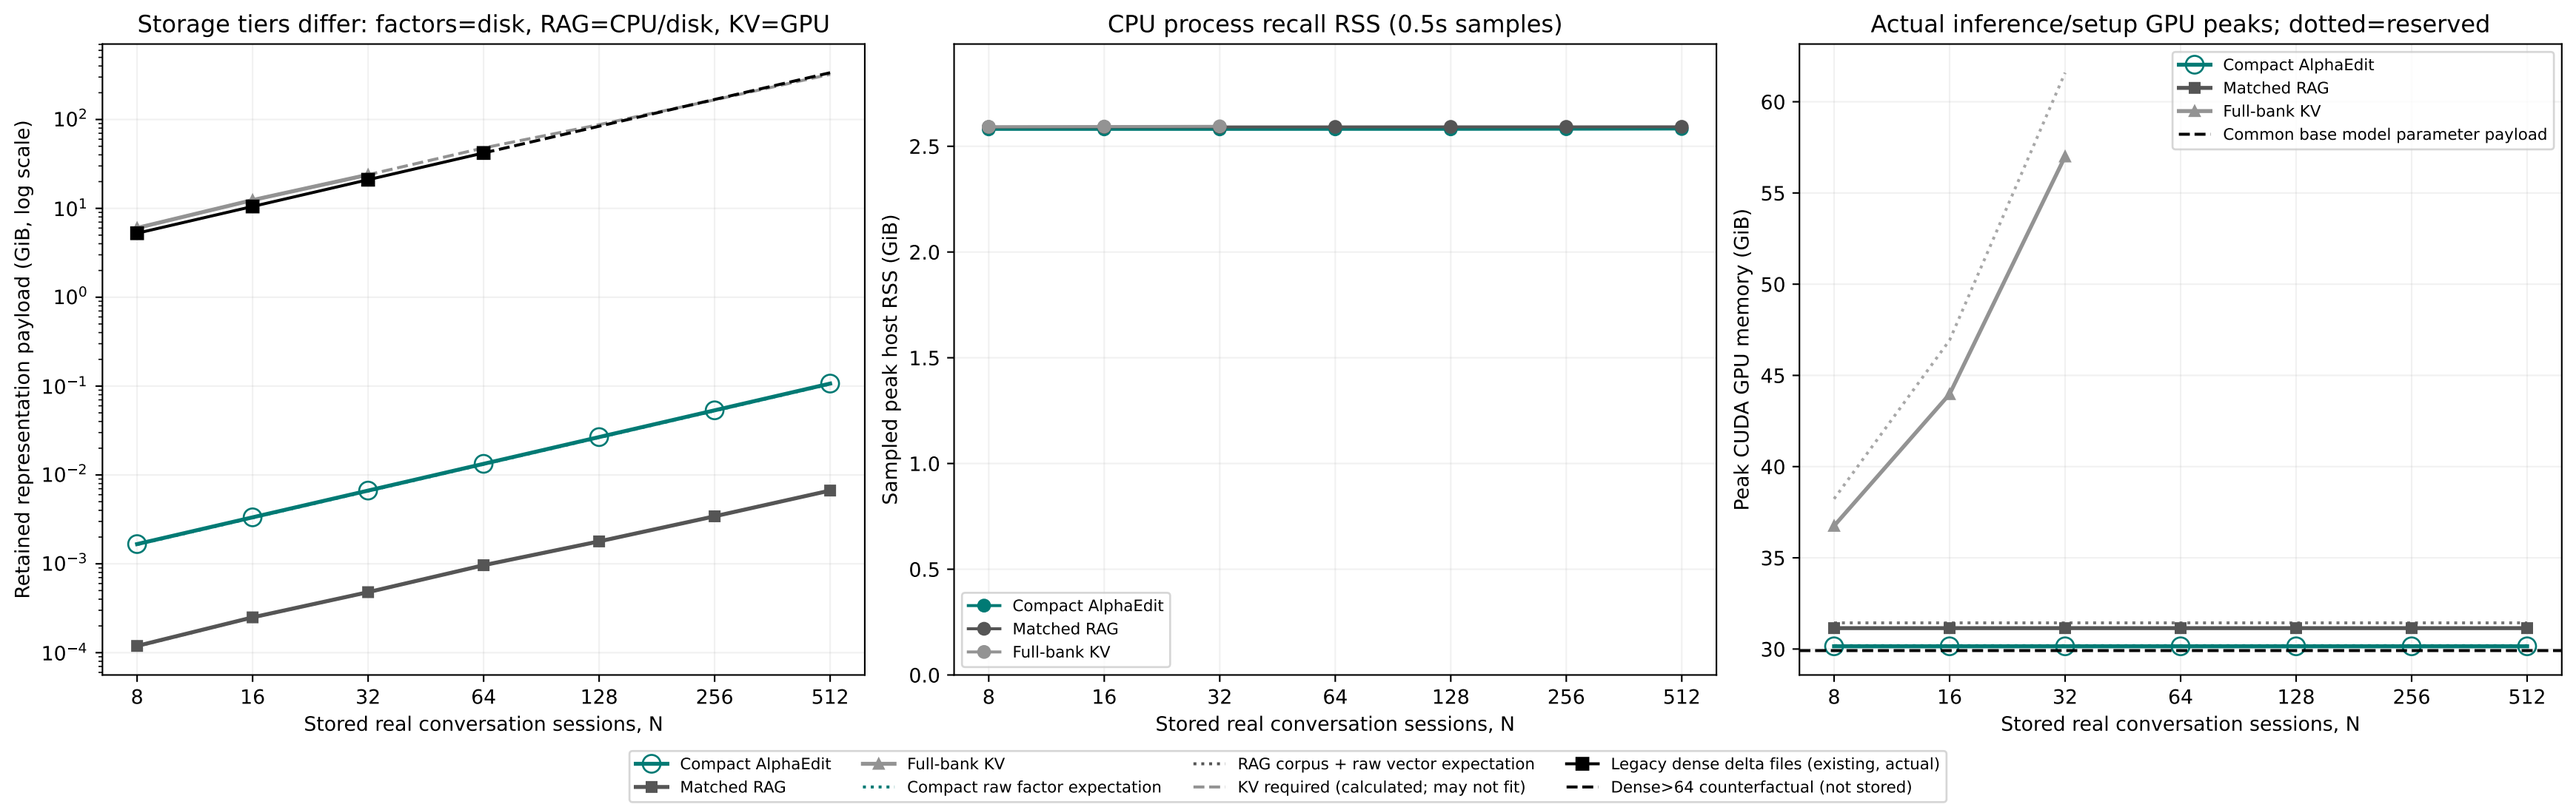
Storage tiers differ: factors=disk, RAG=CPU/disk, KV=GPU
Retained representation payload (GiB, log scale)
Stored real conversation sessions, N
CPU process recall RSS (0.5s samples)
Sampled peak host RSS (GiB)
Stored real conversation sessions, N
Compact AlphaEdit
Matched RAG
Full-bank KV
Actual inference/setup GPU peaks; dotted=reserved
Peak CUDA GPU memory (GiB)
Stored real conversation sessions, N
Compact AlphaEdit
Matched RAG
Full-bank KV
Common base model parameter payload
Compact AlphaEdit
Matched RAG
Full-bank KV
Compact raw factor expectation
RAG corpus + raw vector expectation
KV required (calculated; may not fit)
Legacy dense delta files (existing, actual)
Dense>64 counterfactual (not stored)